

Basics of OpenMP (Recap)



Dirk Schmidl
IT Center, RWTH Aachen University
Member of the HPC Group
schmidl@itc.rwth-aachen.de



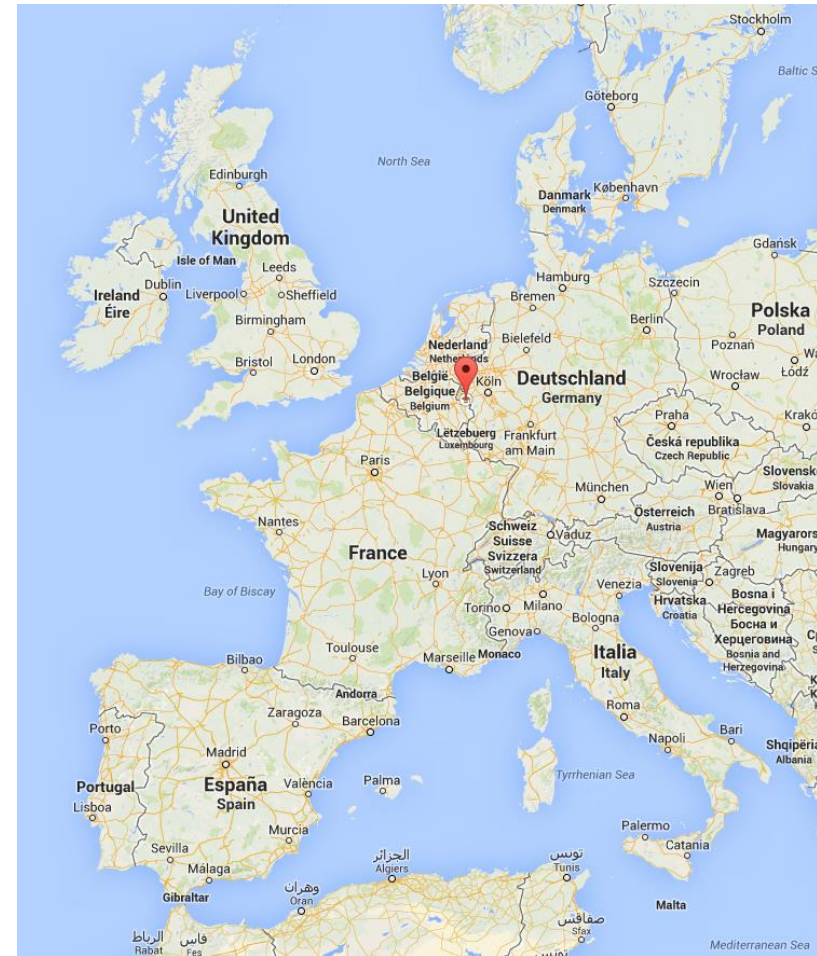
Christian Terboven
IT Center, RWTH Aachen University
Deputy lead of the HPC Group
terboven@itc.rwth-aachen.de

■ RWTH Aachen University

- One of the largest technical universities in Germany.
- ~ 36.000 students
- ~ 500 professors
- ~ 4500 academic staff

■ High Performance Computing Group at the IT Center

- application support /optimization
- focus on shared memory programming
- member of OpenMP ARB



- De-facto standard for Shared-Memory Parallelization.
- 1997: OpenMP 1.0 for FORTRAN
- 1998: OpenMP 1.0 for C and C++
- 1999: OpenMP 1.1 for FORTRAN (errata)
- 2000: OpenMP 2.0 for FORTRAN
- 2002: OpenMP 2.0 for C and C++
- 2005: OpenMP 2.5 now includes both programming languages.
- 05/2008: OpenMP 3.0 release
- 07/2011: OpenMP 3.1 release
- 07/2013: OpenMP 4.0 release



RWTH Aachen University is a member of the OpenMP Architecture Review Board (ARB) since 2006.

Shared Memory Architectures

Single Processor System (dying out)

■ CPU is fast

→ Order of 3.0 GHz

■ Caches:

→ Fast, but expensive


→ Thus small, order of MB

■ Memory is slow

→ Order of 0.3 GHz

→ Large, order of GB


■ A good utilization of caches is crucial for good performance of HPC applications!



Why is there no 4.0 GHz x86 CPU?




- Because that beast would get too hot!



Fast clock cycles make processor chips more expensive, hotter and more power consuming.

Moore's Law still holds!



The number of transistors on a chip is still doubling every 24 months ...

... but the clock speed is no longer increasing that fast!


Instead, we will see many more cores per chip!

Source: Herb Sutter

www.gotw.ca/publications/concurrency-ddj.htm

Dual-Core Processor System

- Since 2005/2006 Intel and AMD are producing dual-core processors for the mass market!
- In 2006/2007 Intel and AMD introduced quad-core processors.
- → Any recently bought PC or laptop is a multi-core system already!




Example for a SMP system

■ Dual-socket Intel Woodcrest (dual-core) system

- Two cores per chip, 3.0 GHz
- Each chip has 4 MB of L2 cache on-chip, shared by both cores
- No off-chip cache
- Bus: Frontsidebus

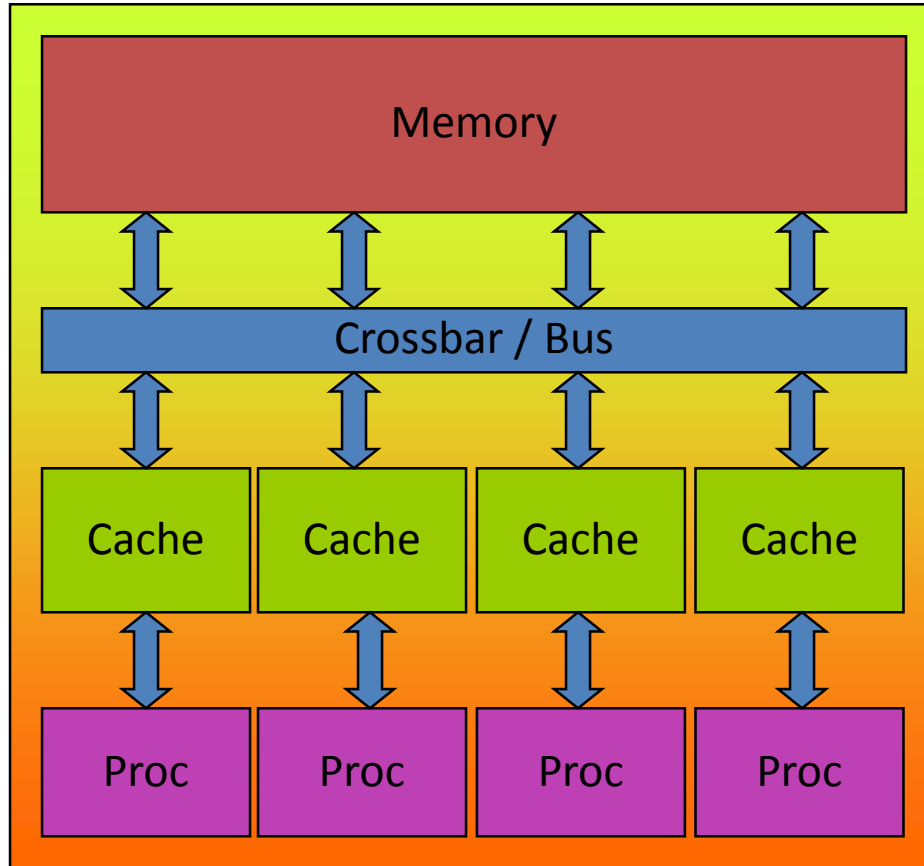
■ SMP: Symmetric Multi Processor

- Memory access time is uniform on all cores
- Limited scalability



OpenMP Overview & Parallel Region & Basic Worksharing

■ OpenMP: Shared-Memory Parallel Programming Model.




All processors/cores access a shared main memory.

Real architectures are more complex, as we will see later / as we have seen.

Parallelization in OpenMP employs multiple threads.

- OpenMP programs start with just one thread: The *Master*.
- *Worker* threads are spawned at *Parallel Regions*, together with the Master they form the *Team* of threads.
- In between Parallel Regions the Worker threads are put to sleep. The OpenMP *Runtime* takes care of all thread management work.
- Concept: *Fork-Join*.
- Allows for an incremental parallelization!



■ The parallelism has to be expressed explicitly.

C/C++

```
#pragma omp parallel
{
    ...
    structured block
    ...
}
```

Fortran

```
!$omp parallel
...
    structured block
...
!$omp end parallel
```

■ **Structured Block**

- Exactly one entry point at the top
- Exactly one exit point at the bottom
- Branching in or out is not allowed
- Terminating the program is allowed
(abort / exit)

■ **Specification of number of threads:**

- ▶ Environment variable:
OMP_NUM_THREADS=...
- ▶ Or: Via `num_threads` clause:
add `num_threads (num)` to the
parallel construct

Hello OpenMP World

Hello orphaned OpenMP World

- From within a shell, global setting of the number of threads:

```
export OMP_NUM_THREADS=4  
./program
```

- From within a shell, one-time setting of the number of threads:

```
OMP_NUM_THREADS=4 ./program
```

- Intel Compiler on Linux: ask the runtime for more information:

```
export KMP_AFFINITY=verbose  
export OMP_NUM_THREADS=4  
./program
```


- If the expression of an `if` clause on a *Parallel Region* evaluates to false
 - The Parallel Region is executed with a Team of one Thread only
 - Used for optimization, e.g. avoid going parallel
- **OpenMP data scoping rules still apply!**

C/C++

```
#pragma omp parallel if(expr)  
...
```

Fortran

```
!$omp parallel if(expr)  
...
```

For Construct

- If only the *parallel* construct is used, each thread executes the Structured Block.
- Program Speedup: *Worksharing*
- OpenMP's most common Worksharing construct: *for*

C/C++

```
int i;  
#pragma omp parallel for  
for (i = 0; i < 100; i++)  
{  
    a[i] = b[i] + c[i];  
}
```

Fortran

```
INTEGER :: i  
!$omp parallel do  
DO i = 0, 99  
    a[i] = b[i] + c[i];  
END DO
```

- Distribution of loop iterations over all threads in a Team.
- Scheduling of the distribution can be influenced.

- Loops often account for most of a program's runtime!

Pseudo-Code
Here: 4 Threads

Serial

```
do i = 0, 99  
  a(i) = b(i) + c(i)  
end do
```

Thread 1

```
do i = 0, 24  
  a(i) = b(i) + c(i)  
end do
```

Thread 2

```
do i = 25, 49  
  a(i) = b(i) + c(i)  
end do
```


Thread 3

```
do i = 50, 74  
  a(i) = b(i) + c(i)  
end do
```

Thread 4

```
do i = 75, 99  
  a(i) = b(i) + c(i)  
end do
```

Memory



Vector Addition

Example: Sparse Matrix Vector Mult.



■ $A = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 2 & 2 & 0 & 0 \\ 0 & 0 & 3 & 0 \\ 4 & 0 & 4 & 4 \end{pmatrix}$

```
for (i = 0; i < num_rows; i++){  
    sum = 0.0;  
    for (nz=row[i]; nz<row[i+1]; ++nz){  
        sum+= value[nz]*x[index[nz]];  
    }  
    y[i] = sum;  
}
```

- Format: compressed row storage
- store all values and columns in arrays (length nnz)
- store beginning of a new row in a third array (length n+1)


value:

1	2	2	3	4	4	4
0	0	1	2	0	2	3

index:

row:

0	1	3	4	7
---	---	---	---	---



Load Imbalance

- **for-construct: OpenMP allows to influence how the iterations are scheduled among the threads of the team, via the *schedule* clause:**
 - ➔ `schedule(static [, chunk])`: Iteration space divided into blocks of chunk size, blocks are assigned to threads in a round-robin fashion. If chunk is not specified: #threads blocks.
 - ➔ `schedule(dynamic [, chunk])`: Iteration space divided into blocks of chunk (not specified: 1) size, blocks are scheduled to threads in the order in which threads finish previous blocks.
 - ➔ `schedule(guided [, chunk])`: Similar to dynamic, but block size starts with implementation-defined value, then is decreased exponentially down to chunk.
- **Default on most implementations is `schedule(static)`.**

Single Construct

C/C++

```
#pragma omp single [clause]  
... structured block ...
```

Fortran

```
!$omp single [clause]  
... structured block ...  
!$omp end single
```

- The `single` construct specifies that the enclosed structured block is executed by only one thread of the team.

→ It is up to the runtime which thread that is.

- Useful for:

→ I/O

→ Memory allocation and deallocation, etc. (in general: setup work)

→ Implementation of the single-creator parallel-executor pattern as we will see now...

Synchronization

■ Can all loops be parallelized with `for`-constructs? No!

→ Simple test: If the results differ when the code is executed backwards, the loop iterations are not independent. BUT: This test alone is not sufficient:

C/C++

```
int i;  
#pragma omp parallel for  
for (i = 0; i < 100; i++)  
{  
    s = s + a[i];  
}
```

■ **Data Race:** If between two synchronization points at least one thread writes to a memory location from which at least one other thread reads, the result is not deterministic (race condition).

- **A *Critical Region* is executed by all threads, but by only one thread simultaneously (*Mutual Exclusion*).**

```
C/C++  
  
#pragma omp critical (name)  
{  
    ... structured block ...  
}
```

- **Do you think this solution scales well?**

```
C/C++  
  
int i;  
#pragma omp parallel for  
for (i = 0; i < 100; i++)  
{  
  
    #pragma omp critical  
    {    s = s + a[i];    }  
  
}
```

It's your turn: Make It Scale!



```
#pragma omp parallel
```

```
{
```

```
#pragma omp for
```

```
for (i = 0; i < 99; i++)
```

```
{
```

```
    s = s + a[i];
```

```
}
```

```
} // end parallel
```

```
do i = 0, 24  
    s = s + a(i)  
end do
```

```
do i = 25, 49  
    s = s + a(i)  
end do
```

```
do i = 0, 99  
    s = s + a(i)  
end do
```



```
do i = 50, 74  
    s = s + a(i)  
end do
```

```
do i = 75, 99  
    s = s + a(i)  
end do
```

- In a *reduction*-operation the operator is applied to all variables in the list. The variables have to be *shared*.

→ `reduction(operator:list)`

→ The result is provided in the associated reduction variable

C/C++

```
#pragma omp parallel for reduction(+:s)
for(i = 0; i < 99; i++)
{
    s = s + a[i];
}
```

→ Possible reduction operators with initialization value:

+ (0), * (1), - (0),

& (~0), | (0), && (1), || (0),

^ (0), min (largest number), max (least number)

■ OpenMP `barrier` (implicit or explicit)

→ All tasks created by any thread of the current *Team* are guaranteed to be completed at barrier exit

```
C/C++
```


```
#pragma omp barrier
```

■ All worksharing constructs have an implied barrier

→ This is a safety net

■ In some cases, the implied barrier can be left out through the “`nowait`” clause

Data Scoping



- All threads have access to the same, globally shared memory
- Data in private memory is only accessible by the thread owning this memory
- No other thread sees the change(s) in private memory
- Data transfer is through shared memory and is 100% transparent to the application

- **Need to get this right**

- Part of the learning curve

- **Private data is undefined on entry and exit**

- Can use firstprivate and lastprivate to address this

- **Each thread has its own temporary view on the data**

- Applicable to shared data only

- Means different threads may temporarily not see the same value for the same variable ...

- All threads have a consistent view of the memory after synchronization constructs

- Technically: synchronization constructs contain a flush construct...

- **Managing the Data Environment is the challenge of OpenMP.**
- **Scoping in OpenMP: Dividing variables in *shared* and *private*:**
 - *private*-list and *shared*-list on Parallel Region
 - *private*-list and *shared*-list on Worksharing constructs
 - General default is *shared* for Parallel Region, *firstprivate* for Tasks.
 - Loop control variables on *for*-constructs are *private*
 - Non-static variables local to Parallel Regions are *private*
 - *private*: A new uninitialized instance is created for each thread
 - *firstprivate*: Initialization with Master's value
 - *lastprivate*: Value of last loop iteration is written back to Master
 - Static variables are *shared*

■ Global / static variables can be privatized with the *threadprivate* directive

- One instance is created for each thread
 - Before the first parallel region is encountered
 - Instance exists until the program ends
 - Does not work (well) with nested Parallel Region
- Based on thread-local storage (TLS)
 - TlsAlloc (Win32-Threads), pthread_key_create (Posix-Threads), keyword `__thread` (GNU extension)

C/C++

```
static int i;  
#pragma omp threadprivate(i)
```

Fortran

```
SAVE INTEGER :: i  
!$omp threadprivate(i)
```

Questions?