

# OpenMP Tasking



Dirk Schmidl  
IT Center, RWTH Aachen University  
Member of the HPC Group  
[schmidl@itc.rwth-aachen.de](mailto:schmidl@itc.rwth-aachen.de)



Christian Terboven  
IT Center, RWTH Aachen University  
Deputy lead of the HPC Group  
[terboven@itc.rwth-aachen.de](mailto:terboven@itc.rwth-aachen.de)

# Tasking

C/C++

```
#pragma omp task [clause]  
... structured block ...
```

Fortran

```
!$omp task [clause]  
... structured block ...  
!$omp end task
```

## ■ Each encountering thread/task creates a new Task

- Code and data is being packaged up
- Tasks can be nested
  - Into another Task directive
  - Into a Worksharing construct

## ■ Data scoping clauses:

- `shared(list)`
- `private(list)`    `firstprivate(list)`
- `default(shared | none)`

- **Some rules from *Parallel Regions* apply:**
  - Static and Global variables are shared
  - Automatic Storage (local) variables are private
  
- **If shared scoping is not derived by default:**
  - Orphaned Task variables are `firstprivate` by default!
  - Non-Orphaned Task variables inherit the `shared` attribute!
  - Variables are `firstprivate` unless `shared` in the enclosing context
  
- **So far no verification tool is available to check Tasking programs for correctness!**

## ■ OpenMP **barrier** (implicit or explicit)

- All tasks created by any thread of the current *Team* are guaranteed to be completed at barrier exit

```
C/C++  
  
#pragma omp barrier
```

## ■ Task barrier: **taskwait**

- Encountering task is suspended until child tasks are complete
- Applies only to direct childs, not descendants!

```
C/C++  
  
#pragma omp taskwait
```

# Example: Fibonacci

# Recursive approach to compute Fibonacci



```
int main(int argc,  
         char* argv[])  
{  
    [...]  
    fib(input);  
    [...]  
}
```

```
int fib(int n)    {  
    if (n < 2) return n;  
    int x = fib(n - 1);  
    int y = fib(n - 2);  
    return x+y;  
}
```

- On the following slides we will discuss three approaches to parallelize this recursive code with Tasking.

# First version parallelized with Tasking (omp-v1)



```
int main(int argc,
        char* argv[])
{
    [...]
    #pragma omp parallel
    {
        #pragma omp single
        {
            fib(input);
        }
    }
    [...]
}
```


```
int fib(int n)    {
    if (n < 2) return n;
    int x, y;
    #pragma omp task shared(x)
    {
        x = fib(n - 1);
    }
    #pragma omp task shared(y)
    {
        y = fib(n - 2);
    }
    #pragma omp taskwait
    return x+y;
}
```

- Only one Task / Thread enters `fib()` from `main()`, it is responsible for creating the two initial work tasks

**Taskwait is required, as otherwise `x` and `y` would be lost**



## ■ Overhead of task creation prevents better scalability!




## ■ Improvement: Don't create yet another task once a certain (small enough) $n$ is reached

```
int main(int argc,
        char* argv[])
{
    [...]
    #pragma omp parallel
    {
        #pragma omp single
        {
            fib(input);
        }
    }
    [...]
}
```

```
int fib(int n)    {
    if (n < 2) return n;
    int x, y;
    #pragma omp task shared(x) \
        if(n > 30)
    {
        x = fib(n - 1);
    }
    #pragma omp task shared(y) \
        if(n > 30)
    {
        y = fib(n - 2);
    }
    #pragma omp taskwait
    return x+y;
}
```

- Speedup is ok, but we still have some overhead when running with 4 or 8 threads

Speedup of Fibonacci with Tasks




- **Improvement: Skip the OpenMP overhead once a certain  $n$  is reached (no issue w/ production compilers)**

```
int main(int argc,
        char* argv[])
{
    [...]
    #pragma omp parallel
    {
        #pragma omp single
        {
            fib(input);
        }
    }
    [...]
}
```

```
int fib(int n)    {
    if (n < 2) return n;
    if (n <= 30)
        return serfib(n);
    int x, y;
    #pragma omp task shared(x)
    {
        x = fib(n - 1);
    }
    #pragma omp task shared(y)
    {
        y = fib(n - 2);
    }
    #pragma omp taskwait
    return x+y;
}
```

■ Everything ok now 😊

Speedup of Fibonacci with Tasks



# Example: Data Scoping

# Data Scoping Example (1/7)



```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;

            // Scope of a:
            // Scope of b:
            // Scope of c:
            // Scope of d:
            // Scope of e:

        }
    }
}
```

# Data Scoping Example (2/7)



```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;

            // Scope of a: shared
            // Scope of b:
            // Scope of c:
            // Scope of d:
            // Scope of e:

        }
    }
}
```



# Data Scoping Example (3/7)



```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;

            // Scope of a: shared
            // Scope of b: firstprivate
            // Scope of c:
            // Scope of d:
            // Scope of e:
        }
    }
}
```

# Data Scoping Example (4/7)



```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;

            // Scope of a: shared
            // Scope of b: firstprivate
            // Scope of c: shared
            // Scope of d:
            // Scope of e:

        }
    }
}
```

# Data Scoping Example (5/7)



```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;

            // Scope of a: shared
            // Scope of b: firstprivate
            // Scope of c: shared
            // Scope of d: firstprivate
            // Scope of e:

        }
    }
}
```

# Data Scoping Example (6/7)



```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;

            // Scope of a: shared
            // Scope of b: firstprivate
            // Scope of c: shared
            // Scope of d: firstprivate
            // Scope of e: private
        }
    }
}
```

Hint: Use default(none) to be forced to think about every variable if you do not see clear.

# Data Scoping Example (7/7)



```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;

            // Scope of a: shared,          value of a: 1
            // Scope of b: firstprivate,    value of b: 0 / undefined
            // Scope of c: shared,          value of c: 3
            // Scope of d: firstprivate,    value of d: 4
            // Scope of e: private,        value of e: 5
        }
    }
}
```

# Task Scheduling

## ■ Task Synchronization explained:

```
#pragma omp parallel num_threads(np)
{
    #pragma omp task
        function_A();
    #pragma omp barrier
    #pragma omp single
    {
        #pragma omp task
            function_B();
    }
}
```

np Tasks created here, one for each thread

All Tasks guaranteed to be completed here

1 Task created here

B-Task guaranteed to be completed here

- **Default: Tasks are *tied* to the thread that first executes them → not necessarily the creator. Scheduling constraints:**
  - Only the thread a task is tied to can execute it
  - A task can only be suspended at task scheduling points
    - Task creation, task finish, `taskwait`, `barrier`, `taskyield`
  - If task is not suspended in a barrier, executing thread can only switch to a direct descendant of all tasks tied to the thread
- **Tasks created with the `untied` clause are never tied**
  - Resume at task scheduling points possibly by different thread
  - ~~No scheduling restrictions, e.g. can be suspended at any point~~
  - But: More freedom to the implementation, e.g. load balancing



- If the expression of an **if** clause on a task evaluates to **false**
  - The encountering task is suspended
  - The new task is executed immediately
  - The parent task resumes when new tasks finishes
  - Used for optimization, e.g. avoid creation of small tasks

- The **taskyield** directive specifies that the current task can be suspended in favor of execution of a different task.

→ Hint to the runtime for optimization and/or deadlock prevention

C/C++

```
#pragma omp taskyield
```

Fortran

```
!$omp taskyield
```

# taskyield Example (1/2)



```
#include <omp.h>

void something_useful();
void something_critical();

void foo(omp_lock_t * lock, int n)
{
    for(int i = 0; i < n; i++)
        #pragma omp task
        {
            something_useful();
            while( !omp_test_lock(lock) ) {
                #pragma omp taskyield
            }
            something_critical();
            omp_unset_lock(lock);
        }
}
```

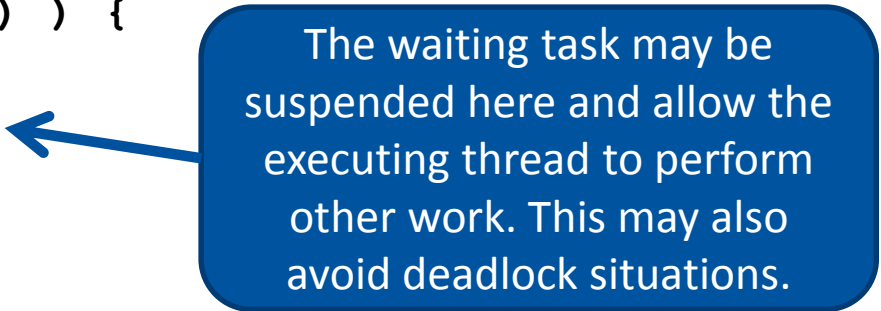
## taskyield Example (2/2)



```
#include <omp.h>

void something_useful();
void something_critical();

void foo(omp_lock_t * lock, int n)
{
    for(int i = 0; i < n; i++)
        #pragma omp task
        {
            something_useful();
            while( !omp_test_lock(lock) ) {
                #pragma omp taskyield
            }
            something_critical();
            omp_unset_lock(lock);
        }
}
```



The waiting task may be suspended here and allow the executing thread to perform other work. This may also avoid deadlock situations.

# Example: Sudoku

## ■ Lets solve Sudoku puzzles with brute multi-core force

	6						8	11			15	14			16
15	11				16	14				12			6		
13		9	12					3	16	14		15	11	10	
2		16		11		15	10	1							
	15	11	10			16	2	13	8	9	12				
12	13			4	1	5	6	2	3					11	10
5		6	1	12		9		15	11	10	7	16			3
	2				10		11	6		5			13		9
10	7	15	11	16				12	13						6
9						1			2		16	10			11
1		4	6	9	13			7		11		3	16		
16	14			7		10	15	4	6	1				13	8
11	10		15				16	9	12	13			1	5	4
		12		1	4	6		16				11	10		
		5		8	12	13		10			11	2			14
3	16			10			7			6				12	

**(1) Find an empty field**

**(2) Insert a number**

**(3) Check Sudoku**

**(4 a) If invalid:**

**Delete number,  
Insert next number**

**(4 b) If valid:**

**Go to next field**

- This parallel algorithm finds all valid solutions

first call contained in a  
`#pragma omp parallel`  
`#pragma omp single`  
 such that one tasks starts the  
 execution of the algorithm

`#pragma omp task`  
 needs to work on a new copy  
 of the Sudoku board

`#pragma omp taskwait`

`#pragma omp taskwait`



first call contained in a

```
#pragma omp parallel
#pragma omp single
```

such that one task starts the  
execution of the algorithm

#pragma omp task  
needs to work on a new copy  
of the Sudoku board

```
#pragma omp taskwait
wait for all child tasks
```

-  (1) Search an empty field
- (2) Insert a number
- (3) Check Sudoku
-  (4 a) If invalid:  
Delete number,  
Insert next number
- (4 b) If valid:  
Go to next field

## ■ OpenMP parallel region creates a team of threads

```
#pragma omp parallel
```

```
{
```

```
#pragma omp single
```

```
    solve_parallel(0, 0, sudoku2, false);
```

```
} // end omp parallel
```

→ Single construct: One thread enters the execution of `solve_parallel`

→ the other threads wait at the end of the `single` ...

→ ... and are ready to pick up threads „from the work queue“

## ■ Syntactic sugar (either you like it or you do not)

```
#pragma omp parallel sections
```

```
{
```

```
    solve_parallel(0, 0, sudoku2, false);
```

```
} // end omp parallel
```



## ■ The actual implementation

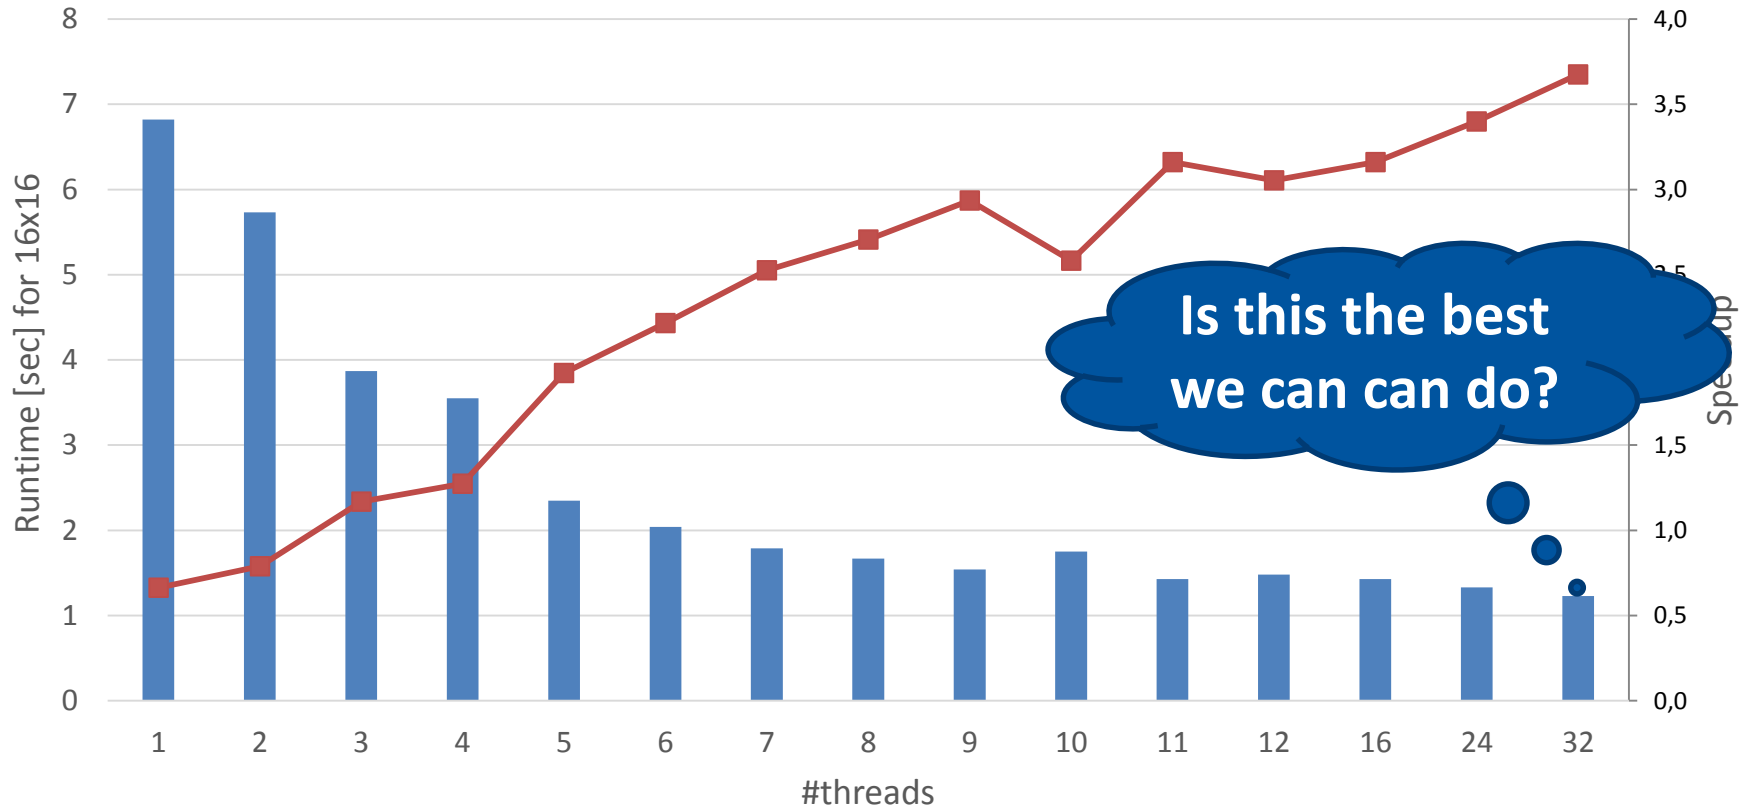
```
for (int i = 1; i <= sudoku->getFieldSize(); i++) {  
    if (!sudoku->check(x, y, i)) {  
#pragma omp task firstprivate(i,x,y,sudoku)  
{  
    // create from copy constructor  
    CSudokuBoard new_sudoku(*sudoku);  
    new_sudoku.set(y, x, i);  
    if (solve_parallel(x+1, y, &new_sudoku)) {  
        new_sudoku.printBoard();  
    }  
} // end omp task  
}  
}
```

#pragma omp task  
needs to work on a new copy  
of the Sudoku board

#pragma omp taskwait  
wait for all child tasks

Sudoku on 2x Intel Xeon E5-2650 @2.0 GHz

Intel C++ 13.1, scatter binding      speedup: Intel C++ 13.1, scatter binding



Is this the best we can do?

# taskgroup Construct

C/C++

```
#pragma omp taskgroup  
... structured block ...
```

Fortran

```
!$omp taskgroup  
... structured block ...  
!$omp end task
```

- **Specifies a wait on completion of child tasks and their descendent tasks**
  - „deeper“ synchronization than `taskwait`, but
  - with the option to restrict to a subset of all tasks (as opposed to a `barrier`)
- **Main use case for now in OpenMP 4.0: Cancellation...**

# Questions?