

Further OpenMP 4.0 Features



Dirk Schmidl
IT Center, RWTH Aachen University
Member of the HPC Group
schmidl@itc.rwth-aachen.de

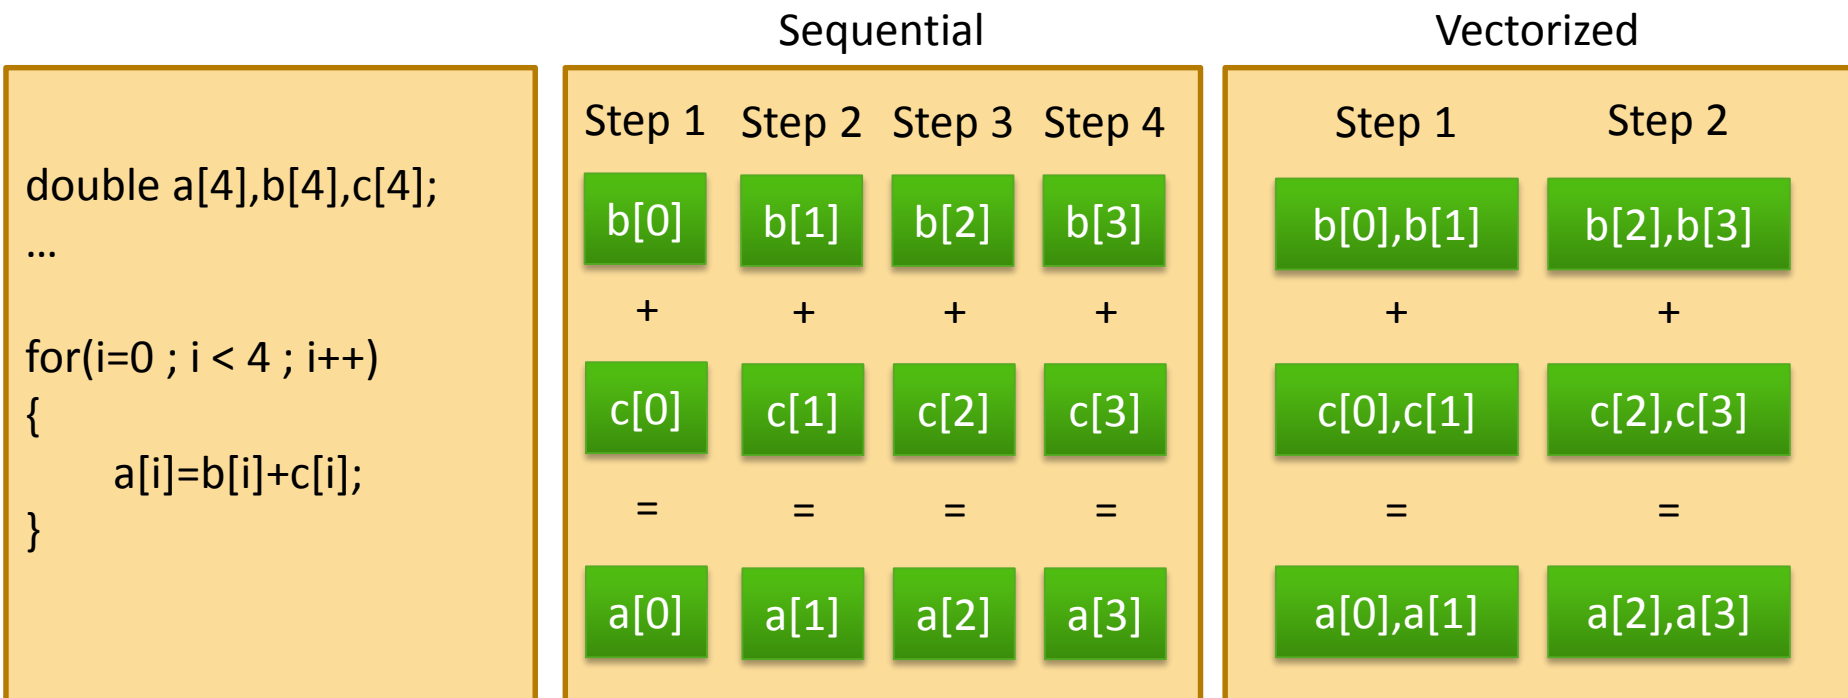


Christian Terboven
IT Center, RWTH Aachen University
Deputy lead of the HPC Group
terboven@itc.rwth-aachen.de

Vectorization

■ SIMD = Single Instruction Multiple Data

- Special hardware instructions to operate on multiple data points at once
- Instructions work on vector registers
- Vector length is hardware dependent



■ Vector lengths on Intel architectures

→ 128 bit: SSE = Streaming SIMD Extensions



→ 256 bit: AVX = Advanced Vector Extensions




→ 512 bit: AVX-512




■ Vectorization works best on aligned data structures.


Good alignment



Bad alignment



Very bad alignment




Compiler
auto-vectorization

Explicit Vector Programming
(e.g. with OpenMP)

Inline Assembly
(e.g.)

Assembler Code
(e.g. addps, mulpd, ...)



The OpenMP SIMD constructs

- The SIMD construct enables the execution of multiple iterations of the associated loops concurrently by means of SIMD instructions.

C/C++:

```
#pragma omp simd [clause(s)]  
for-loops
```

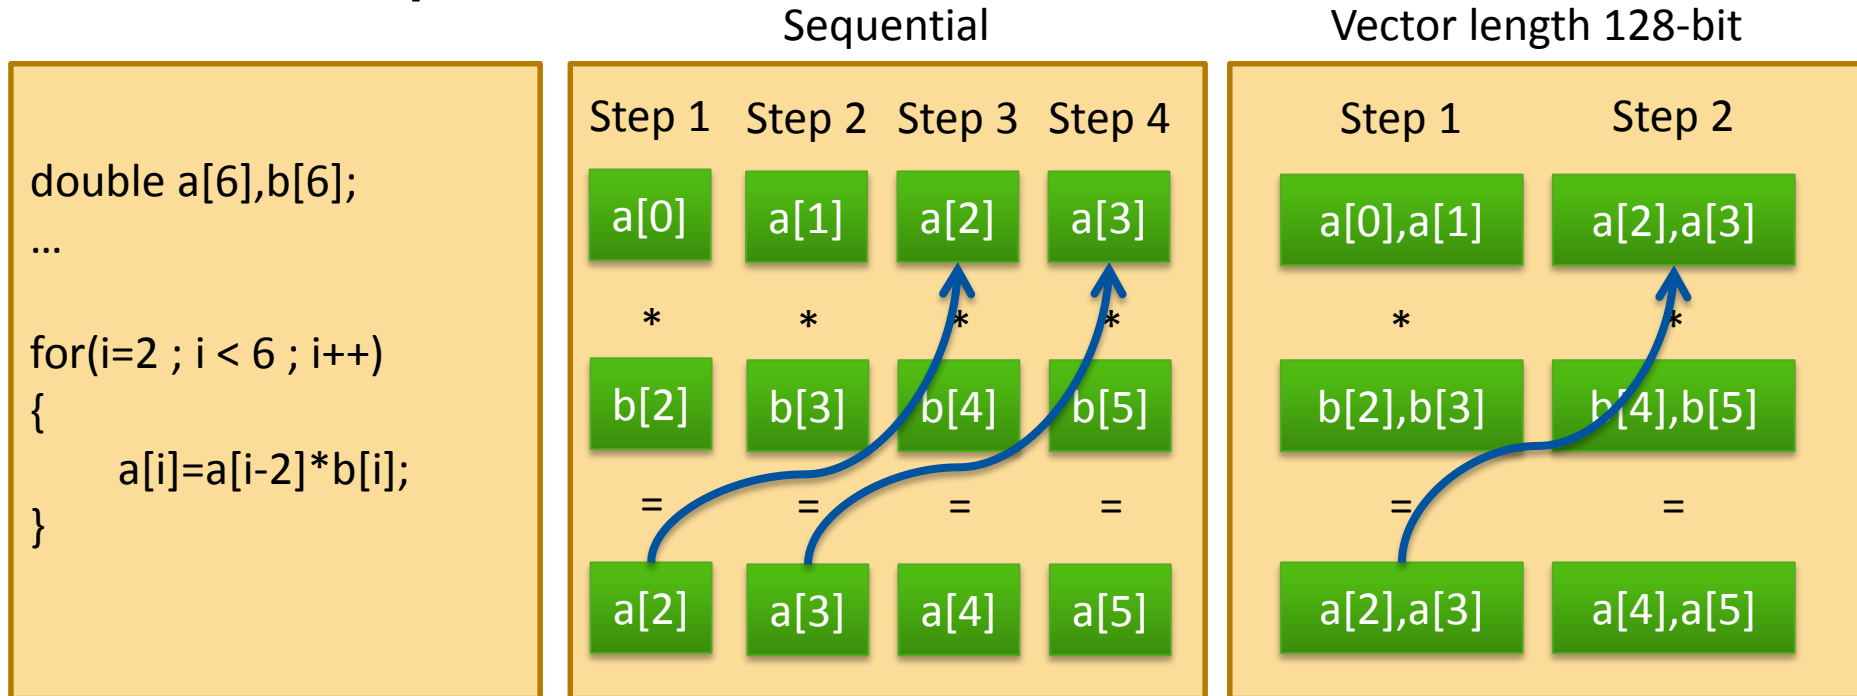
Fortran:

```
!$omp simd [clause(s)]  
do-loops  
[!$omp end simd]
```


- where clauses are:

- linear(*list[:linear-step]*), a variable increases linearly in every loop iteration
- aligned(*list[:alignment]*), specifies that data is aligned
- private(*list*), as usual
- lastprivate(*list*) , as usual
- reduction(*reduction-identifier:list*) , as usual
- collapse(*n*), collapse loops first, and then apply SIMD instructions

- The safelen clause allows to specify a distance of loop iterations where no dependencies occur.



- The safelen clause allows to specify a distance of loop iterations where no dependencies occur.



- Any vector length smaller than or equal to the length specified by safelen can be chosen for vectorization.
- In contrast to parallel for/do loops the iterations are executed in a specified order.

- The loop SIMD construct specifies a loop that can be executed in parallel by all threads and in SIMD fashion on each thread.

C/C++:

```
#pragma omp for simd [clause(s)]  
for-loops
```


Fortran:

```
!$omp do simd [clause(s)]  
do-loops  
[!$omp end do simd [nowait]]
```

- Loop iterations are first distributed across threads, then each chunk is handled as a SIMD loop.
- Clauses:
 - All clauses from the *loop*- or SIMD-construct are allowed
 - Clauses which are allowed for both constructs are applied twice, once for the threads and once for the SIMDization.

- **Function calls in SIMD-loops can lead to bottlenecks, because functions need to be executed serially.**

```
for(i=0 ; i < N ; i++)  
{  
    a[i]=b[i]+c[i];  
    d[i]=sin(a[i]);  
    e[i]=5*d[i];  
}
```



Solutions:

- avoid or inline functions
- create functions which work on vectors instead of scalars

- Enables the creation of multiple versions of a function or subroutine where one or more versions can process multiple arguments using SIMD instructions.

C/C++:

```
#pragma omp declare simd [clause(s)]  
[#pragma omp declare simd [clause(s)]]  
    function definition / declaration
```

Fortran:

```
!$omp declare simd (proc_name)[clause(s)]
```

- where clauses are:

- `simdlen(length)`, the number of arguments to process simultaneously
- `linear(list[:linear-step])`, a variable increases linearly in every loop iteration
- `aligned(argument-list[:alignment])`, specifies that data is aligned
- `uniform(argument-list)`, arguments have an invariant value
- `inbranch / notinbranch`, function is always/never called from within a conditional statement

PI Example Code

File: f.c


```
#pragma omp declare simd
double f(double x)
{
    return (4.0 / (1.0 + x*x));
}
```

File: pi.c

```
#pragma omp declare simd
double f(double x);
...
#pragma omp simd linear(i) private(fX) reduction(+:fSum)
for (i = 0; i < n; i++)
{
    fX = fH * ((double)i + 0.5);
    fSum += f(fX);
}
return fH * fSum;
```

Calculating Pi with
numerical integration
of:

$$\pi = \int_0^1 \frac{4}{1+x^2}$$



■ Runtime of the benchmark on:

- Westmere CPU with SSE (128-bit vectors)
- Intel Xeon Phi with AVX-512 (512-bit vectors)


	Runtime Westmere	Speedup Westmere	Runtime Xeon Phi	Speedup Xeon Phi
non vectorized	1.44 sec	1	16.25 sec	1
vectorized	0.72 sec	2	1.82 sec	8.9

Note: Speedup for memory bound applications might be lower on both systems.

Cancellation

Searching in a binary tree data structure


- Tasks can be created for the left and the right subtree, leading to a task parallel search.
- Multiple threads can pick up the tasks and execute them.
- The search can be finished when one thread found the desired value.



How to end the search process if one thread found the value?

- A shared variable can be used which is set when the value is found.
- All threads check for this variable at every node and stop when the variable is set.
- A lot of flushes are needed to make sure the variable is synchronized.
- All tasks need to be dequeued to check the variable.

=> This is a very inelegant and error prone solution.



- Users can request cancellation of a construct
- threads / tasks will be canceled and execution continues after the end of the construct

C/C++:

```
#pragma omp cancel [construct-type]
```

Fortran:

```
!$omp cancel [construct-type]
```

- Types are: parallel, do/for, taskgroup, sections
- threads/ tasks stop execution at a certain point and not immediately
- Users can add cancellation points manually:

C/C++:

```
#pragma omp cancellation point
```

Fortran:

```
!$omp cancellation point
```

Task Dependencies

C/C++

```
#pragma omp task depend(dependency-type: list)
... structured block ...
```

- The *task dependence* is fulfilled when the predecessor task has completed
 - `in` dependency-type: the generated task will be a dependent task of all previously generated sibling tasks that reference at least one of the list items in an `out` or `inout` clause.
 - `out` and `inout` dependency-type: The generated task will be a dependent task of all previously generated sibling tasks that reference at least one of the list items in an `in`, `out`, or `inout` clause.
 - The list items in a `depend` clause may include array sections.

Degree of parallelism exploitable in this concrete example:
T2 and T3 (2 tasks), T1 of next iteration has to wait for them

T1 has to be completed before T2 and T3 can be executed.

T2 and T3 can be executed in parallel.

```
void process_in_parallel() {  
    #pragma omp parallel  
    #pragma omp single  
    {  
        int x = 1;  
        ...  
        for (int i = 0; i < T; ++i) {  
            #pragma omp task shared(x, ...) depend(out: x) // T1  
            preprocess_some_data(...);  
            #pragma omp task shared(x, ...) depend(in: x) // T2  
            do_something_with_data(...);  
            #pragma omp task shared(x, ...) depend(in: x) // T3  
            do_something_independent_with_data(...);  
        }  
    } // end omp single, omp parallel  
}
```

- The following allows for more parallelism, as there is one *i* per thread. Hence, two tasks may be active per thread.

```
void process_in_parallel() {  
    #pragma omp parallel  
    {  
        #pragma omp for  
        for (int i = 0; i < T; ++i) {  
            #pragma omp task depend(out: i)  
                preprocess_some_data(...);  
            #pragma omp task depend(in: i)  
                do_something_with_data(...);  
            #pragma omp task depend(in: i)  
                do_something_independent_with_data(...);  
        }  
    } // end omp parallel  
}
```


- The following allows for even more parallelism, as there now can be two tasks active per thread per i-th iteration.

```
void process_in_parallel() {  
    #pragma omp parallel  
    #pragma omp single  
    {  
        for (int i = 0; i < T; ++i) {  
            #pragma omp task firstprivate(i)  
            {  
                #pragma omp task depend(out: i)  
                preprocess_some_data(...);  
                #pragma omp task depend(in: i)  
                do_something_with_data(...);  
                #pragma omp task depend(in: i)  
                do_something_independent_with_data(...);  
            } // end omp task  
        }  
    } // end omp single, end omp parallel  
}
```


„Real“ Task Dependencies



```
void blocked_cholesky( int NB, float A[NB][NB] ) {
    int i, j, k;
    for (k=0; k<NB; k++) {
        #pragma omp task depend(inout:A[k][k])
        spotrf (A[k][k]) ;
        for (i=k+1; i<NT; i++)
            #pragma omp task depend(in:A[k][k]) depend(inout:A[k][i])
            strsm (A[k][k], A[k][i]);
        // update trailing submatrix
        for (i=k+1; i<NT; i++) {
            for (j=k+1; j<i; j++)
                #pragma omp task depend(in:A[k][i],A[k][j])
                    depend(inout:A[j][i])
                sgemm( A[k][i], A[k][j], A[j][i]);
            #pragma omp task depend(in:A[k][i]) depend(inout:A[i][i])
            ssyrk (A[k][i], A[i][i]);
        }
    }
}
```



* image from BSC

User Defined Reductions

Atomics

- Use `declare reduction` directive to define operators
- Operators used in reduction clause like predefined ops

```
#pragma omp declare reduction (reduction-identifier :  
typename-list : combiner) [initializer(initializer-expr)]
```

- **reduction-identifier** gives a name to the operator
 - Can be overloaded for different types
 - Can be redefined in inner scopes
- **typename-list** is a list of types to which it applies
- **combiner expression** specifies how to combine values
- **initializer** specifies the operator's identity value
 - `initializer-expression` is an expression or a function

■ Declare the reduction operator

```
#pragma omp declare reduction (merge : std::vector<int> :  
    omp_out.insert(omp_out.end(), omp_in.begin(), omp_in.end()))
```

■ Use the reduction operator in a reduction clause

```
void schedule (std::vector<int> &v, std::vector<int> &filtered) {  
    #pragma omp parallel for reduction (merge : filtered)  
    for (std::vector<int>::iterator it = v.begin(); it < v.end();  
        it++)  
        if ( filter(*it) )    filtered.push_back(*it);  
}
```

■ Private copies created for a reduction are initialized to the identity that was specified for the operator and type

→ Default identity defined if `identity` clause not present

■ Compiler uses `combiner` to combine private copies

→ `omp_out` refers to private copy that holds combined value

→ `omp_in` refers to the other private copy

■ Use `atomic` construct for mutually exclusive access to a single memory location

```
#pragma omp atomic [read|write|update|capture]  
[seq_cst]  
expression-stmt
```

- `expression-stmt` restricted based on type of `atomic`
- `update`, the default behavior, reads and writes the single memory location atomically
- `read` reads location atomically
- `write` writes location atomically
- `capture` updates or writes location and captures its value (before or after update) into a private variable

Questions?